

SIMATS ENGINEERING
ASSESSMENT TOOL 1: Analytical Problem Solving

**CSA1511 – CLOUD COMPUTING AND BIG DATA
ANALYTICS**

Name : NAVADEEP R

Reg No : 192424110

Faculty Incharge : Dr.HEMAVATHI R

CO2 – ASSESSMENT TOOL – 1

Analytical Problem Solving

1. Compare Type-1 and Type-2 hypervisors for a cloud datacenter hosting mission-critical applications. Analyze performance, security, and manageability, then justify the better choice.

Problem Statement

A cloud datacenter needs to host mission-critical applications that require high performance, strong security, and easy management. The two common hypervisor types are Type-1 (Bare-Metal) and Type-2 (Hosted) hypervisors.

Analysis

Factor	Type-1 Hypervisor	Type-2 Hypervisor
Architecture	Runs directly on physical hardware	Runs on top of a host operating system
Performance	Very high	Lower due to OS overhead
Security	Strong isolation and smaller attack surface	More dependent on host OS security
Resource Management	Efficient allocation of CPU, RAM, storage	Additional overhead from host OS
Manageability	Designed for centralized datacenter management	Easier for desktop/testing environments

Factor	Type-1 Hypervisor	Type-2 Hypervisor
Examples	VMware ESXi, Microsoft Hyper-V, Xen	VirtualBox, VMware Workstation

Performance

- Type-1 hypervisors communicate directly with the physical hardware. Therefore, they provide lower overhead and better CPU, memory, storage, and network performance.
- Type-2 hypervisors require a host operating system between the hypervisor and hardware. This introduces additional processing overhead.

Security

- Type-1 hypervisors provide stronger workload isolation because they do not depend on a general-purpose host operating system.
- Type-2 hypervisors have a larger potential attack surface because vulnerabilities in the host OS can affect the virtualized environment.

Manageability

Type-1 hypervisors are designed for datacenter environments and support centralized VM management, resource allocation, monitoring, migration, and high availability.

Justification

For mission-critical cloud applications, Type-1 hypervisors are the better choice because they provide:

1. Higher performance and lower virtualization overhead.
2. Better security and VM isolation.
3. Efficient resource utilization.
4. Better centralized management and scalability.

Conclusion

A Type-1 hypervisor is more suitable for cloud datacenters because mission-critical workloads require high availability, performance, security, and efficient resource management.

2. A cloud provider must isolate workloads belonging to finance, healthcare, and student portals on the same hardware. Analyze how virtualization architecture supports isolation and identify two major attack surfaces.

Problem Statement

A cloud provider needs to run finance, healthcare, and student portal workloads on the same physical hardware. Since these workloads contain sensitive information, they must remain isolated from one another.

How Virtualization Provides Isolation

Virtualization divides a physical server into multiple Virtual Machines (VMs). Each VM has its own virtual CPU, memory, storage, operating system, and applications.

For example:

Physical Server

- VM 1 – Finance Applications
- VM 2 – Healthcare Applications
- VM 3 – Student Portal

The hypervisor controls access to physical resources and prevents one VM from directly accessing another VM's resources.

Important Isolation Mechanisms

1. Memory Isolation

Each VM receives a separate virtual memory space. One VM cannot directly access another VM's memory.

2. CPU Isolation

CPU resources are scheduled between VMs, preventing one workload from directly controlling another workload's CPU execution.

3. Storage Isolation

Separate virtual disks and access permissions prevent unauthorized access to other workloads' files.

4. Network Isolation

VLANs, virtual switches, firewalls, and security groups can separate finance, healthcare, and student traffic.

Two Major Attack Surfaces

1. Hypervisor Attack Surface

The hypervisor is a critical component because it controls multiple VMs. A vulnerability in the hypervisor could potentially allow an attacker to escape from a VM and access the host or other VMs.

This is known as a VM Escape Attack.

2. Virtual Network Attack Surface

Virtual switches, virtual network interfaces, and network configurations can be attacked through techniques such as unauthorized access, traffic interception, or network misconfiguration.

Conclusion

Virtualization provides strong isolation through separate CPU, memory, storage, and network resources. However, the hypervisor and virtual networking layers must be secured because they are major attack surfaces.

3. Study the following VM host utilization snapshot and recommend corrective actions: CPU 92%, memory 88%, storage I/O latency 35 ms, average VM boot time 170 s. Explain the likely bottlenecks.

Problem Statement

The following VM host utilization data is provided:

- CPU utilization: 92%
- Memory utilization: 88%
- Storage I/O latency: 35 ms
- Average VM boot time: 170 seconds

The objective is to identify the bottlenecks and recommend corrective actions.

Analysis of the Metrics

Parameter	Current Value	Analysis
CPU	92%	Very high utilization
Memory	88%	High memory pressure
Storage I/O latency	35 ms	High I/O delay
VM boot time	170 s	Slow VM startup

1. CPU Bottleneck

CPU utilization of 92% indicates that the host is heavily loaded.

Possible causes:

- Too many VMs running on the host.
- CPU-intensive applications.
- Insufficient CPU allocation.
- Poor workload distribution.

Corrective action:

Move some VMs to another host using VM migration/load balancing, or increase available CPU resources.

2. Memory Bottleneck

Memory usage of 88% indicates that RAM resources are approaching their capacity.

When memory becomes insufficient, the system may use swapping or memory reclamation, which can reduce performance.

Corrective action:

- Add more physical RAM.
- Reduce unnecessary VM memory allocation.
- Move some VMs to another host.
- Optimize memory-intensive applications.

3. Storage I/O Bottleneck

Storage I/O latency of 35 ms indicates that disk operations are taking significant time.

Possible causes include:

- High number of simultaneous disk operations.
- Slow storage devices.
- Storage contention between VMs.
- Insufficient IOPS.

Corrective action:

- Use faster SSD/NVMe storage.
- Distribute workloads across storage devices.
- Increase available IOPS.
- Optimize disk-intensive applications.

4. VM Boot-Time Problem

An average boot time of 170 seconds is relatively high.

The likely reasons include:

- CPU and memory contention.
- Slow storage I/O.
- Large VM images.
- Heavy startup applications.

Recommended Overall Solution

The cloud provider should:

1. Load-balance VMs across multiple hosts.
2. Increase CPU and RAM capacity.
3. Upgrade storage to faster SSD/NVMe infrastructure.
4. Monitor resource usage continuously.
5. Use automated VM migration and resource scheduling.

Conclusion

The main bottlenecks are high CPU utilization, high memory utilization, and storage I/O latency. These resource constraints are likely contributing to the slow VM boot time of 170 seconds. Resource balancing, hardware upgrades, and workload optimization can improve overall VM performance.

4. Analyze how Software Defined Datacenter architecture improves agility compared with a traditional datacenter. Support your answer with compute, storage, and network virtualization considerations.

Problem Statement

A traditional datacenter generally depends heavily on manually configured physical servers, storage devices, and network equipment.

A Software Defined Datacenter (SDDC) uses software-defined technologies to virtualize and centrally manage compute, storage, and networking resources.

Traditional Datacenter

Physical Servers → Physical Storage → Physical Network Devices

Management is often performed separately for each infrastructure component.

Software Defined Datacenter

Software Management Layer

↓

Virtualized Compute + Virtualized Storage + Virtualized Network

This enables centralized and automated infrastructure management.

Compute Virtualization

SDDC uses virtual machines to divide physical servers into multiple logical computing environments.

Benefits:

- Faster VM provisioning.
- Better hardware utilization.
- Easy workload migration.
- Dynamic resource allocation.

Storage Virtualization

Storage virtualization combines physical storage resources and presents them as logical storage pools.

Benefits:

- Centralized storage management.
- Better storage utilization.
- Easier scaling.
- Automated storage allocation.

Network Virtualization

Network virtualization creates logical networks independent of physical network hardware.

Technologies such as virtual switches, software-defined networking, and network policies can be used.

Benefits:

- Faster network configuration.
- Network segmentation.
- Automated security policies.
- Flexible traffic management.

How SDDC Improves Agility

Traditional Datacenter	Software Defined Datacenter
Manual configuration	Software-based automation
Physical resource dependency	Virtualized resources
Slow provisioning	Rapid provisioning
Separate management	Centralized management
Difficult scaling	Easy scaling
Limited flexibility	Highly flexible

Example

Suppose an organization needs 50 new VMs.

In a traditional environment, administrators may need to configure physical servers, networking, and storage manually.

In an SDDC, software-based management can automatically provision the required VMs and allocate the required compute, storage, and networking resources.

Conclusion

SDDC improves datacenter agility by transforming compute, storage, and networking into software-controlled, virtualized resources. This enables faster provisioning, automation, scalability, centralized management, and efficient resource utilization.

5. A multi-cloud federation project fails due to identity mismatches between providers. Analyze the issue and propose a secure identity and privacy design using federation concepts.

Problem Statement

A company uses multiple cloud providers. However, the federation project fails because each provider uses different identity systems, user identifiers, authentication mechanisms, and authorization policies.

For example:

Cloud Provider A

User ID: user123

Cloud Provider B

User ID: student456

Even though both identities may represent the same person, the providers cannot automatically recognize them as the same identity.

Causes of Identity Mismatch

1. Different identity formats.
2. Different authentication mechanisms.
3. Different authorization models.
4. Different user attributes and roles.
5. Lack of a common trust relationship.

Proposed Secure Identity Federation

A centralized Identity Provider (IdP) can be introduced.

Example architecture:

User



Identity Provider (IdP)

↓ Authentication

Federation / Trust Layer



Cloud Provider A Cloud Provider B

The Identity Provider authenticates the user once and provides trusted identity information to the participating cloud providers.

Federation Technologies

Protocols such as SAML 2.0 and OpenID Connect (OIDC) can be used for identity federation.

The identity provider can issue security assertions or tokens containing information such as:

- User identity
- Role
- Organization

- Permissions
- Authentication status

Secure Identity Design

1. Single Sign-On (SSO)

Users authenticate through a trusted identity provider rather than maintaining separate passwords for each cloud provider.

2. Role-Based Access Control

Roles can be mapped between providers.

Example:

University Role: Student

↓

Cloud A Role: Student-User

↓

Cloud B Role: Limited-User

This ensures that users receive only the permissions they require.

3. Multi-Factor Authentication

MFA should be enabled at the identity provider to provide an additional security layer.

4. Token-Based Authentication

Short-lived tokens should be used instead of sharing passwords between cloud providers.

5. Privacy Protection

Only the minimum required user information should be shared with each provider.

For example, if Cloud B only needs the user's role, the federation system should avoid sending unnecessary personal information.

6. Trust Relationship

Cloud providers should establish trusted relationships with the Identity Provider and validate signed authentication assertions/tokens.

Privacy-Aware Federation Flow

User Login



Identity Provider authenticates user



MFA verification



Identity/token generated



Cloud Provider validates token



Role and permissions mapped



User receives authorized cloud service

Conclusion

The identity mismatch can be solved through identity federation, where a trusted Identity Provider provides a common authentication mechanism across multiple cloud providers. Using SAML/OIDC, SSO, MFA, role mapping, short-lived tokens, and minimum data sharing provides both secure identity management and privacy protection.